

# Using Trusted Execution Environments for Secure Stream Processing of Medical Data

DAIS'19 - 18-20 July 2019, Copenhagen (DK) - *Use-Case Paper*

**C. Segarra**<sup>1</sup> R. Delgado<sup>1</sup> M. Lemay<sup>1</sup> P. Aublin<sup>3</sup> P. Pietzuch<sup>3</sup> V. Schiavoni<sup>2</sup>

<sup>1</sup>CSEM, Neuchâtel, Switzerland, {carlos.segarra,ricard.delgado,mathieu.lemay}@csem.ch

<sup>2</sup>University of Neuchâtel, Switzerland, valerio.schiavoni@unine.ch

<sup>3</sup>Imperial College London, United Kingdom, {p.aublin,prp}@imperial.ac.uk

Thursday, June 13, 2019

## ① Introduction

Motivation

## ② Background

Technical Background

Medical Use-Case

## ③ Secure Stream Processing of Medical Data

Project Goals

Architecture

Evaluation & Results

## ④ Conclusions and Further Work



# Motivation

Why do we need secure big data processing engines?

**We want to:**

**Process** large amounts of **sensitive** information



# Motivation

Why do we need secure big data processing engines?

**We want to:**

**Process** large amounts of **sensitive** information



Generally requires...

**Outsourcing** of data **storage** and **processing**



# Motivation

Why do we need secure big data processing engines?

**We want to:**

**Process** large amounts of **sensitive** information

↓ Generally requires...

**Outsourcing** of data **storage** and **processing**

↓ Generally means...

Cloud tenant **can** access **protected** data



# Motivation

Why do we need secure big data processing engines?

**We want to:**

**Process** large amounts of **sensitive** information

↓ Generally requires...

**Outsourcing** of data **storage** and **processing**

↓ Generally means...

Cloud tenant **can** access **protected** data

↓ This is...

**UNACCEPTABLE FOR INDUSTRIES IN THE MEDICAL DOMAIN**  
**(and many others)**



# Trusted Execution Environments (TEE)

Formal definition, examples and availability

## Trusted Execution Environment

A **TEE** is a secure area of a processor. It guarantees code and data to be protected with respect to **confidentiality** and **integrity**.



# Trusted Execution Environments (TEE)

Formal definition, examples and availability

## Trusted Execution Environment

A **TEE** is a secure area of a processor. It guarantees code and data to be protected with respect to **confidentiality** and **integrity**.

- Available on a variety of commodity CPUs. For instance:





# Trusted Execution Environments (TEE)

Formal definition, examples and availability

## Trusted Execution Environment

A **TEE** is a secure area of a processor. It guarantees code and data to be protected with respect to **confidentiality** and **integrity**.



**Intel SGX** Available on consumer-grade CPUs starting from architecture codename *Skylake*.

**arm**  
TRUSTZONE

**Arm Trustzone** Available on Cortex-A processors and v8 Cortex-M23 and Cortex-M33 (e.g. Raspberry Pi).



# Intel Software Guard eXtensions

## Definition, Threat Model and Known Vulnerabilities

- **Intel Software Guard eXtensions (SGX)** are a set of instructions and memory access extensions that enable applications to create **hardware-protected areas** in their address space called **enclaves**.
- Security perimeter includes **only** the internals of the CPU package.
- An **attestation protocol** verifies that code is running in a **genuine enclave** and that it has not been tampered.




Figure 1: Intel SGX Operating Principles.



# Intel Software Guard eXtensions

## Definition, Threat Model and Known Vulnerabilities

- **Intel Software Guard eXtensions (SGX)** are a set of instructions and memory access extensions that enable applications to create **hardware-protected areas** in their address space called **enclaves**.
- Security perimeter includes **only** the internals of the CPU package.
- An **attestation protocol** verifies that code is running in a **genuine enclave** and that it has not been tampered.




Figure 1: Intel SGX Operating Principles.



# Intel Software Guard eXtensions

## Definition, Threat Model and Known Vulnerabilities

- Intel Software Guard eXtensions (SGX) are a set of instructions and memory access extensions that enable applications to create **hardware-protected areas** in their address space called **enclaves**.
- Security perimeter includes **only** the internals of the CPU package.
- An **attestation protocol** verifies that code is running in a **genuine enclave** and that it has not been tampered.

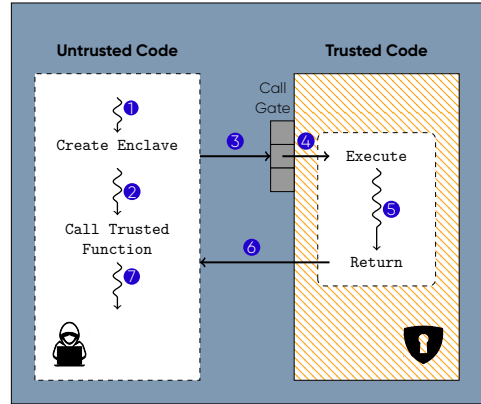


Figure 1: Intel SGX Operating Principles.



# SGX-Spark

## Running Spark Jobs inside Enclaves

### SGX-Spark

- SGX-Spark is a framework under-development at the **Imperial College London** to enable seamless deployment of **Spark** jobs inside **enclaves**.
- Protect **confidentiality** and **integrity** of existing Spark jobs **without** modifications to the application code.
- Execute only **sensitive** parts of the application **inside** the enclave. Leave information outside the enclave **encrypted**.

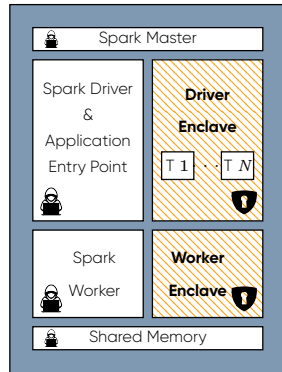


Figure 2: SGX-Spark Architecture.



# SGX-Spark

## Running Spark Jobs inside Enclaves

### SGX-Spark

- SGX-Spark is a framework under-development at the **Imperial College London** to enable seamless deployment of **Spark** jobs inside **enclaves**.
- Protect **confidentiality** and **integrity** of existing Spark jobs **without** modifications to the application code.
- Execute only **sensitive** parts of the application **inside** the enclave. Leave information outside the enclave **encrypted**.




Figure 2: SGX-Spark Architecture.



# SGX-Spark

## Running Spark Jobs inside Enclaves

### SGX-Spark

- SGX-Spark is a framework under-development at the **Imperial College London** to enable seamless deployment of **Spark** jobs inside **enclaves**.
- Protect **confidentiality** and **integrity** of existing Spark jobs **without** modifications to the application code.
- Execute only **sensitive** parts of the application **inside** the enclave. Leave information outside the enclave **encrypted**.

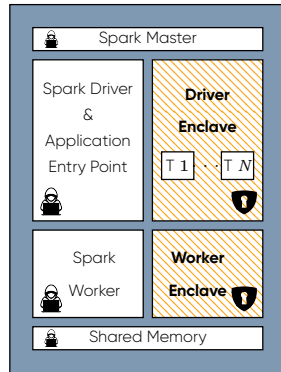


Figure 2: SGX-Spark Architecture.



# Medical Use-Case

## Cardiac Monitoring, ECG, PPG, and HRV

- Our solution is **independent of the chosen data streams**, hence it could be leveraged in a variety of domains where privacy might be a concern. However,
- The data streams used for this project belong to the **medical** domain, in particular, they are obtained from human **cardiac activity monitoring**. The two most standard procedures are:
  - ① **ECG**: measure heart's electrical activity over time.  
*E.g: chest straps.*
  - ② **PPG**: measure blood's volume variation over time.  
*E.g: smartwatches and pulse oximeters.*





# Medical Use-Case

## Cardiac Monitoring, ECG, PPG, and HRV

- Our solution is **independent of the chosen data streams**, hence it could be leveraged in a variety of domains where privacy might be a concern. However,
- The data streams used for this project belong to the **medical** domain, in particular, they are obtained from human **cardiac activity monitoring**. The two most standard procedures are:
  - ① **ECG**: measure heart's electrical activity over time.  
*E.g:* chest straps.
  - ② **PPG**: measure blood's volume variation over time.  
*E.g:* smartwatches and pulse oximeters.



# Medical Use-Case

## Cardiac Monitoring, ECG, PPG, and HRV

- Our solution is **independent of the chosen data streams**, hence it could be leveraged in a variety of domains where privacy might be a concern. However,
- The data streams used for this project belong to the **medical** domain, in particular, they are obtained from human **cardiac activity monitoring**. The two most standard procedures are:
  - ① **ECG**: measure heart's electrical activity over time.  
*E.g:* chest straps.
  - ② **PPG**: measure blood's volume variation over time.  
*E.g:* smartwatches and pulse oximeters.



# Medical Use-Case

## Cardiac Monitoring, ECG, PPG, and HRV

- We are interested in the **inter-beat intervals** from the generated diagram (23 – 69 B per second per user).
- From an **ECG** we can obtain obtain these intervals from the time between **R peaks**, abbreviated as RR intervals (Figure 3).
- With their **timestamps** we compute a live analysis of the **Heart Rate Variability (HRV)**.




Figure 3: Schematic representation of an ECG signal and the data we gather.



# Medical Use-Case

## Cardiac Monitoring, ECG, PPG, and HRV

- We are interested in the **inter-beat intervals** from the generated diagram (23 – 69 B per second per user).
- From an **ECG** we can obtain obtain these intervals from the time between **R peaks**, abbreviated as RR intervals (Figure 3).
- With their **timestamps** we compute a live analysis of the **Heart Rate Variability (HRV)**.




Figure 3: Schematic representation of an ECG signal and the data we gather.



# Medical Use-Case

## Cardiac Monitoring, ECG, PPG, and HRV

- We are interested in the **inter-beat intervals** from the generated diagram (23 – 69 B per second per user).
- From an **ECG** we can obtain obtain these intervals from the time between **R peaks**, abbreviated as RR intervals (Figure 3).
- With their **timestamps** we compute a live analysis of the **Heart Rate Variability (HRV)**.




Figure 3: Schematic representation of an ECG signal and the data we gather.



# Secure Stream Processing of Medical Data

## Project Goals

### Project Goals

- Use SGX-Spark to implement a **streaming platform** that gathers data from sensors and **securely** outsources computation.
  - Deploy in an existing medical environment.
  - Perform stress tests to assess the system's robustness and reliability.
  - Evaluate and quantify the overhead of providing strong security guarantees.



# Secure Stream Processing of Medical Data

## Envisioned Scenario

The **chosen scenario** involves:

- ① ECG data streamed from a **sensor** to a **gateway**
- ② Real-time processing with **HRV** algorithms: SDNN and HRV Bands analysis.
  - **SDNN**: rolling standard deviation of NN (RR) intervals.
  - **HRV Bands**: rolling Discrete Fourier Transform and low/high frequency component.
- ③ Support for **storage** and result **post-processing**




Figure 4: Wearable-enabled ecosystem of our platform.



# System Architecture

## Component Analysis and Execution Workflow




Figure 5: Client-server architecture.





# System Architecture

## Component Analysis and Execution Workflow




Figure 5: Client-server architecture.



# System Architecture

## Component Analysis and Execution Workflow




Figure 5: Client-server architecture.



# System Architecture

## Component Analysis and Execution Workflow




Figure 5: Client-server architecture.



# System Architecture

## Component Analysis and Execution Workflow



### Execution Flow


- 1 **Sensors** generate samples and stream them to the **gateway** over **mqtt**
- 2 **Gateways** aggregate data, encrypts it and sends it to the **server's** filesystem interface over **SFTP**

Figure 5: Client-server architecture.



# System Architecture

## Component Analysis and Execution Workflow



### Execution Flow


- 1 **Sensors** generate samples and streams them to the **gateway** over **mqtt**
- 2 **Gateways** aggregate data, encrypts it and sends it to the **server's** filesystem interface over **SFTP**
- 3 **SGX-Spark** streaming job monitors the FS, **batch-processes** data, and saves encrypted results back in the FS

Figure 5: Client-server architecture.



# System Architecture

## Component Analysis and Execution Workflow



### Execution Flow

- 1 **Sensors** generate samples and streams them to the **gateway** over **mqtt**
- 2 **Gateways** aggregate data, encrypts it and sends it to the **server's** filesystem interface over **SFTP**
- 3 **SGX-Spark** streaming job monitors the FS, **batch-processes** data, and saves encrypted results back in the FS
- 4 **Gateways** fetch result files

Figure 5: Client-server architecture.



# Evaluation & Results

## Experimental Setup & Metrics

### Three Execution Modes:

- 1 Vanilla Spark
- 2 SGX-Spark w/o Enclaves
- 3 **SGX-Spark w/ Enclaves**

### Two Algorithms:

- 1 Identity (Batch & Stream)
- 2 SDNN (Batch & Stream)

### Two Metrics:



- 1 Elapsed Time
- 2 Avg. Batch Processing Time

Workload	s_rate (samples / sec)	Input Load
Batch - Small	{44, 89, 178, 356, 712, 1424}	{1, 2, 4, 8, 16, 32} kB
Stream - Small	{44, 89, 178, 356, 712, 1424}	{1, 2, 4, 8, 16, 32} kB / sec
Batch - Big	{44, 89, 178, 356, 712, 1424} * 1024	{1, 2, 4, 8, 16, 32} MB
Stream - Big	{44, 89, 178, 356, 712, 1424} * 1024	{1, 2, 4, 8, 16, 32} MB / sec



# Evaluation & Results

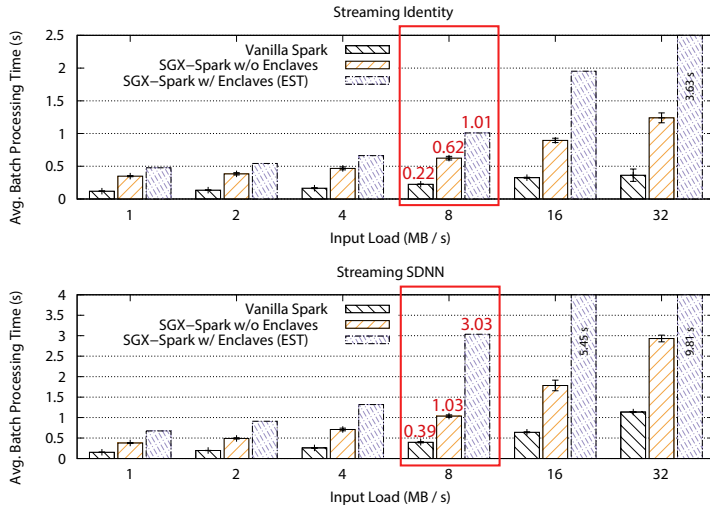
## Results: Batch Execution - Big Load





# Evaluation & Results

## Results: Stream Execution – Big Load



# Conclusion

- Introduced a PoC of a **privacy-preserving streaming platform**.
  - Introduces  $4 \times - 5 \times$  slowdown vs. **vanilla Spark Streaming** (load  $< 4$  MB per second).
  - Requires **no changes** to the application source code.
- Further lines of research:
  - Perform an economical evaluation of the cost of deploying our system to the cloud:  
**how expensive is privacy?**
  - **Reduce the TCB** on the client side leveraging TEEs for low-power devices (e.g. ARM TrustZone).



# Conclusion

- Introduced a PoC of a **privacy-preserving streaming platform**.
  - Introduces  $4 \times - 5 \times$  slowdown vs. **vanilla Spark Streaming** (load  $< 4$  MB per second).
  - Requires **no changes** to the application source code.
- Further lines of research:
  - Perform an economical evaluation of the cost of deploying our system to the cloud:  
**how expensive is privacy?**
  - **Reduce the TCB** on the client side leveraging TEEs for low-power devices (e.g. ARM TrustZone).

**Thank you very much for your attention! Questions?**

